

## Deploying on Kubernetes- Jupyter Notebook

1. Create a Kubernetes cluster on GCP - `gcloud container clusters create --machine-type n1-standard-2 --num-nodes 2 --zone us-central1-a --cluster-version latest adityadwkubernetescluster`
2. `docker pull adityadw/jupyter_notebook`
3. `docker tag adityadw/jupyter_notebook:latest gcr.io/genuine-grid-327615/adityadw/jupyter_notebook:1`
4. `docker push gcr.io/genuine-grid-327615/adityadw/jupyter_notebook:1`

```
adidwivedi96@cloudshell:~ (genuine-grid-327615)$ docker pull adityadw/jupyter_notebook
Using default tag: latest
latest: Pulling from adityadw/jupyter_notebook
284055322776: Pull complete
5031b1992b2a: Pull complete
b6329deb9320: Pull complete
b6769b3e91cb: Pull complete
c3b5faa0dcfb: Pull complete
Digest: sha256:758ab9230acb3707902c3bc1c2604b711dd570c24eee6925dcc85a4028582b81
Status: Downloaded newer image for adityadw/jupyter_notebook:latest
docker.io/adityadw/jupyter_notebook:latest
adidwivedi96@cloudshell:~ (genuine-grid-327615)$ docker tag adityadw/jupyter_notebook:latest gcr.io/genuine-grid-327615/adityadw/jupyter_notebook:1
adidwivedi96@cloudshell:~ (genuine-grid-327615)$ docker push gcr.io/genuine-grid-327615/adityadw/jupyter_notebook:1
The push refers to repository [gcr.io/genuine-grid-327615/adityadw/jupyter_notebook]
38301074c5ed: Pushed
aa81112ca16f: Pushed
2889f2c4cf15: Pushed
2d652f97bd28: Pushed
824bf068fd3d: Layer already exists
1: digest: sha256:758ab9230acb3707902c3bc1c2604b711dd570c24eee6925dcc85a4028582b81 size: 1371
```

5. Create deployment file for jupyter notebook as follows, here the internal port on which service is running is 8888 -

```
jupyter_notebook_deployment.yaml > ...
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: jupyter
5    labels:
6      app: jupyter
7  spec:
8    selector:
9      matchLabels:
10       app: jupyter
11    replicas: 2
12    minReadySeconds: 15
13    strategy:
14      type: RollingUpdate
15      rollingUpdate:
16        maxUnavailable: 1
17        maxSurge: 1
18    template:
19      metadata:
20        labels:
21          app: jupyter
22      spec:
23        containers:
24          - image: adityadw/jupyter_notebook
25            imagePullPolicy: Always
26            name: jupyter
27            ports:
28              - containerPort: 8888
29
```

6. `kubectl apply -f jupyter_notebook_deployment.yaml`
7. Create service file in order to deploy jupyter notebook as follows. Presently it has been set as loadbalancing service so that the testing can be done via browser -

```

jupyter_service.yaml x jupyter_notebook_deployment.yaml
jupyter_service.yaml > ...
1  apiVersion: v1
2  kind: Service
3  metadata:
4    name: jupyter
5  spec:
6    type: LoadBalancer
7    ports:
8      - port: 8888
9        protocol: TCP
10       targetPort: 8888
11    selector:
12      app: jupyter

```

## 8. kubectl apply -f jupyter\_service.yaml

```

adidwivedi96@cloudshell:~/clud-infra/Project (genuine-grid-327615)$ kubectl apply -f jupyter_notebook_deployment.yaml
deployment.apps/jupyter created
adidwivedi96@cloudshell:~/clud-infra/Project (genuine-grid-327615)$ kubectl apply -f jupyter_service.yaml
service/jupyter created
adidwivedi96@cloudshell:~/clud-infra/Project (genuine-grid-327615)$ kubectl apply -f jupyter_service.yaml
service/jupyter configured
adidwivedi96@cloudshell:~/clud-infra/Project (genuine-grid-327615)$

```

## 9. We have the service running on public IP as follows -

Name	Last Modified	File size
AzureSearchIndex - Met Artworks.ipynb	12 days ago	5.91 kB
Classification - Adult Census with Vowpal Wabbit.ipynb	12 days ago	5.07 kB
Classification - Adult Census.ipynb	12 days ago	3.34 kB
Classification - Before and After MMLSpark.ipynb	12 days ago	10.2 kB
Classification - Twitter Sentiment with Vowpal Wabbit.ipynb	12 days ago	11.7 kB
CognitiveServices - Celebrity Quote Analysis.ipynb	12 days ago	7.59 kB
CognitiveServices - Overview.ipynb	12 days ago	34.8 kB
CognitiveServices - Predictive Maintenance.ipynb	12 days ago	15.8 kB
ConditionalKNN - Exploring Art Across Cultures.ipynb	12 days ago	12 kB
CyberML - Anomalous Access Detection.ipynb	12 days ago	15.4 kB
DeepLearning - BiLSTM Medical Entity Extraction.ipynb	12 days ago	11.6 kB
DeepLearning - CIFAR10 Convolutional Network.ipynb	12 days ago	4.34 kB
DeepLearning - Flower Image Classification.ipynb	12 days ago	6.01 kB
DeepLearning - Transfer Learning.ipynb	12 days ago	3.72 kB
HttpOnSpark - Working with Arbitrary Web APIs.ipynb	12 days ago	2.54 kB
HyperParameterTuning - Fighting Breast Cancer.ipynb	12 days ago	4.58 kB
Interpretability - Image Explainers.ipynb	12 days ago	12.2 kB
Interpretability - Tabular SHAP explainer.ipynb	12 days ago	10.4 kB
Interpretability - Text Explainers.ipynb	12 days ago	8.17 kB

## Apache Spark

1. Follow the steps as above to load the docker image from the third party to gcr -  
docker pull bitnami/spark
2. docker tag bitnami/spark:latest gcr.io/genuine-grid-327615/bitnami/spark:1
3. docker push gcr.io/genuine-grid-327615/bitnami/spark:1

```

didwivedi96@cloudshell:~/cloud-infra/Project (genuine-grid-327615)$ docker pull bitnami/spark
Using default tag: latest
latest: Pulling from bitnami/spark
3b664af592e: Pull complete
da89ddc65be: Pull complete
d21c74b4228: Pull complete
0c51956a4b: Pull complete
b28f6fabedd: Pull complete
51828bf66a9: Pull complete
cd8e55f83fb: Pull complete
7800b61fa67: Pull complete
89ae2b97f94: Pull complete
78a7ee58f2a: Pull complete
Digest: sha256:373048f025775d4f8b093f1fa704e87f8603469e5bfe5068c18855c00affca1
Status: Downloaded newer image for bitnami/spark:latest
docker.io/bitnami/spark:latest
didwivedi96@cloudshell:~/cloud-infra/Project (genuine-grid-327615)$ docker tag bitnami/spark:latest gcr.io/genuine-grid-327615/bitnami/spark:1
didwivedi96@cloudshell:~/cloud-infra/Project (genuine-grid-327615)$ docker push gcr.io/genuine-grid-327615/bitnami/spark:1
The push refers to repository [gcr.io/genuine-grid-327615/bitnami/spark]
572bf3e81bc: Pushed
6ef99e19e35: Pushed
31f0eaf1d69: Pushed
2cc2969a3da: Pushed
a31da56e68e: Pushed
d9bb1c6eaffa: Pushed
f5d1f7c253a: Pushed
353863c9e7e: Pushed
e5896b4e039: Pushed
75a52abf3b1: Layer already exists

```

#### 4. Create deployment file as follows -

```

spark_deployment.yaml > ...
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: spark
5    labels:
6      app: spark
7  spec:
8    selector:
9      matchLabels:
10       app: spark
11    replicas: 2
12    minReadySeconds: 15
13    strategy:
14      type: RollingUpdate
15      rollingUpdate:
16        maxUnavailable: 1
17        maxSurge: 1
18    template:
19      metadata:
20        labels:
21          app: spark
22      spec:
23        containers:
24          - image: bitnami/spark
25            imagePullPolicy: Always
26            name: spark
27          ports:
28            - containerPort: 8080
29

```

#### 5. kubectl apply -f spark\_deployment.yaml

6. Create service file as follows, again we are using loadbalancer service so that we get a public IP which can be tested in the browser

```
spark_service.yaml > ...
1  apiVersion: v1
2  kind: Service
3  metadata:
4    name: spark
5  spec:
6    type: LoadBalancer
7    ports:
8      - port: 8080
9        protocol: TCP
10     targetPort: 8080
11  selector:
12    app: spark
```

```
1: digest: sha256:373048f025775d4f8b093f1fa704e87f8603469e5bfe5068c18855c00affca1 size: 2425
adidwivedi96@cloudshell:~/clud-infra/Project (genuine-grid-327615) $ kubectl apply -f spark_deployment.yaml
deployment.apps/spark configured
adidwivedi96@cloudshell:~/clud-infra/Project (genuine-grid-327615) $ kubectl apply -f spark_service.yaml
service/spark unchanged
```

7. We see the following output on the public IP

Spark Master at spark://spark-f9bd9c64f-8x7lj:7077

URL: spark://spark-f9bd9c64f-8x7lj:7077  
Alive Workers: 0  
Cores in use: 0 Total, 0 Used  
Memory in use: 0.0 B Total, 0.0 B Used  
Resources in use:  
Applications: 0 Running, 0 Completed  
Drivers: 0 Running, 0 Completed  
Status: ALIVE

Workers (0)

Worker Id	Address	State	Cores	Memory	Resources
-----------	---------	-------	-------	--------	-----------

Running Applications (0)

Application ID	Name	Cores	Memory per Executor	Resources Per Executor	Submitted Time	User	State	Duration
----------------	------	-------	---------------------	------------------------	----------------	------	-------	----------

Completed Applications (0)

Application ID	Name	Cores	Memory per Executor	Resources Per Executor	Submitted Time	User	State	Duration
----------------	------	-------	---------------------	------------------------	----------------	------	-------	----------

## Apache Hadoop

1. docker pull bde2020/hadoop-namenode:2.0.0-hadoop3.2.1-java8
2. docker tag bde2020/hadoop-namenode:2.0.0-hadoop3.2.1-java8 gcr.io/genuine-grid-327615/bde2020/hadoop-namenode:1
3. docker push gcr.io/genuine-grid-327615/bde2020/hadoop-namenode:1

4. `docker pull bde2020/hadoop-datanode:2.0.0-hadoop3.2.1-java8`
5. `docker tag bde2020/hadoop-datanode:2.0.0-hadoop3.2.1-java8 gcr.io/genuine-grid-327615/bde2020/hadoop-datanode:1`
6. `docker push gcr.io/genuine-grid-327615/bde2020/hadoop-datanode:1`

```
7948c3e5790c: Layer already exists
12 digest: sha256:51ad9293ec52083e5003ef0aabb00c3dd7d6335dd4f495ec1257f97a272cbb4c0 size: 3033
adidwiedi96@cloudshell:~/cloud-infra/Project (genuine-grid-327615) $ docker pull bde2020/hadoop-datanode:2.0.0-hadoop3.2.1-java8
2.0.0-hadoop3.2.1-java8: Pulling from bde2020/hadoop-datanode
3192219afd04: Already exists
7127a1d8cced: Already exists
883a89599900: Already exists
77920a3e82af: Already exists
92329e91aecd: Already exists
f373218fec59: Already exists
aa53513fec97: Already exists
8b1800105b98: Already exists
c3a84a3e49c8: Already exists
a65640a64a76: Already exists
3ca2ec07878c: Pull complete
26c2dd45430e: Pull complete
13c9c87a46cb: Pull complete
Digest: sha256:ddfe9ad55af4f73d2ccb6da31d9e331fffb94d5f046126db4f40aa348d484bf
Status: Downloaded newer image for bde2020/hadoop-datanode:2.0.0-hadoop3.2.1-java8
docker.io/bde2020/hadoop-datanode:2.0.0-hadoop3.2.1-java8
adidwiedi96@cloudshell:~/cloud-infra/Project (genuine-grid-327615) $ docker tag bde2020/hadoop-datanode:2.0.0-hadoop3.2.1-java8 gcr.io/genuine-grid-327615/bde2020/hadoop-datanode:1
adidwiedi96@cloudshell:~/cloud-infra/Project (genuine-grid-327615) $ docker push gcr.io/genuine-grid-327615/bde2020/hadoop-datanode:1
The push refers to repository [gcr.io/genuine-grid-327615/bde2020/hadoop-datanode]
6a6a1867fd39: Pushed
70da79e16472: Pushed
```

7. Create deployment yaml for master apache hadoop node as follows. Here we have env section which sets our hadoop cluster's name. The internal port to be exposed is 9870 -

hadoop\_master.yaml > {} spec > {} strategy > {} rollingUpdate > ...

```
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: master-hadoop
5    labels:
6      app: master-hadoop
7  spec:
8    selector:
9      matchLabels:
10       app: master-hadoop
11    replicas: 2
12    minReadySeconds: 15
13    strategy:
14      type: RollingUpdate
15      rollingUpdate:
16        maxUnavailable: 1
17        maxSurge: 1
18    template:
19      metadata:
20        labels:
21          app: master-hadoop
22      spec:
23        containers:
24          - image: bde2020/hadoop-namenode
25            imagePullPolicy: Always
26            name: master-hadoop
27            env:
28              - name: CLUSTER_NAME
29                value: "cloud-infra-hadoop"

    ports:
      - containerPort: 9870
```

(tabbed space not shown)

8. `kubectl apply -f hadoop_master.yaml`
9. Create the service file for hadoop master as follows. Here again we have a loadbalancer service which maps internal port 9870 to external 9870 -

```

hadoop_master_service.yaml > ...
2   kind: Service
3   metadata:
4     name: master-hadoop
5   spec:
6     type: LoadBalancer
7     ports:
8     - port: 9870
9       protocol: TCP
10      targetPort: 9870
11    selector:
12      app: master-hadoop

```

#### 10. kubectl apply -f hadoop\_master\_service.yaml

```

addiwivedi96@cloudshell:~/cloud-infra/Project (genuine-grid-327615)$ kubectl apply -f hadoop_master.yaml
deployment.apps/master-hadoop created
addiwivedi96@cloudshell:~/cloud-infra/Project (genuine-grid-327615)$ kubectl apply -f hadoop_master_service.yaml
service/hadoop-master created

```

#### 11. Create worker hadoop deployment as follows. Here there is an environment variable which points the hadoop worker to master using kubedns.

```

hadoop_worker.yaml > {} spec > {} selector > ...
1   apiVersion: apps/v1
2   kind: Deployment
3   metadata:
4     name: worker-hadoop
5     labels:
6       app: worker-hadoop
7   spec:
8     selector:
9       matchLabels:
10        app: worker-hadoop
11    replicas: 2
12    minReadySeconds: 15
13    strategy:
14      type: RollingUpdate
15      rollingUpdate:
16        maxUnavailable: 1
17        maxSurge: 1
18    template:
19      metadata:
20        labels:
21          app: worker-hadoop
22      spec:
23        containers:
24        - image: bde2020/hadoop-datanode
25          imagePullPolicy: Always
26          name: worker-hadoop
27          env:
28          - name: SERVICE_PRECONDITION
29            value: "http://master-hadoop:9870"

```

#### 11. Replicate the same deployment file for 2nd hadoop worker as well with the change in name and labels as follows -

```

hadoop_worker_2.yaml > {} spec > {} template > {} spec > containers > ...
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: worker-hadoop-2
5    labels:
6      app: worker-hadoop-2
7  spec:
8    selector:
9      matchLabels:
10       app: worker-hadoop-2
11    replicas: 2
12    minReadySeconds: 15
13    strategy:
14      type: RollingUpdate
15      rollingUpdate:
16        maxUnavailable: 1
17        maxSurge: 1
18    template:
19      metadata:
20        labels:
21          app: worker-hadoop-2
22      spec:
23        containers:
24          - image: bde2020/hadoop-datanode
25            imagePullPolicy: Always
26            name: worker-hadoop-2
27            env:
28              - name: SERVICE_PRECONDITION
29                value: "http://master-hadoop:9870"

```

12. `kubectl apply -f hadoop_worker.yaml`


13. `kubectl apply -f hadoop_worker_2.yaml`

```

deployment.apps/worker-hadoop configured
adidwivedi96@cloudshell:~/cloud-infra/Project (genuine-grid-327615)$ kubectl apply -f hadoop_worker.yaml
deployment.apps/worker-hadoop configured
adidwivedi96@cloudshell:~/cloud-infra/Project (genuine-grid-327615)$ kubectl apply -f hadoop_worker_2.yaml
deployment.apps/worker-hadoop-2 created
adidwivedi96@cloudshell:~/cloud-infra/Project (genuine-grid-327615)$

```

14. We get the following hadoop master view on its public IP



## Overview 'master-hadoop-c97f44696-c8hxt:8020' (active)

Started:	Thu Oct 28 02:54:17 -0400 2021
Version:	3.2.1, rb3cbbb467e22ea829b3808f4b7b01d07e0bf3842
Compiled:	Tue Sep 10 11:56:00 -0400 2019 by rohithsharmaks from branch-3.2.1
Cluster ID:	CID-a21fc68c-b0b6-40f7-8fb2-235a808f9eb8
Block Pool ID:	BP-237167288-10.8.1.11-1635404055091

## Summary

Security is off.

Safemode is off.

1 files and directories, 0 blocks (0 replicated blocks, 0 erasure coded block groups) = 1 total filesystem object(s).

Heap Memory used 66.66 MB of 114.13 MB Heap Memory. Max Heap Memory is 1.76 GB.

Non Heap Memory used 40.88 MB of 42.56 MB Committed Non Heap Memory. Max Non Heap Memory is <unbounded>.

## Sonarqube and Sonar-Scanner

1. `docker pull adityadw/sonarqube_sonarscanner`

2. `docker tag adityadw/sonarqube_sonarscanner:latest gcr.io/genuine-grid-327615/adityadw/sonarqube_sonarscanner:1`
3. `docker push gcr.io/genuine-grid-327615/adityadw/sonarqube_sonarscanner:1`

```
adidwivedi96@cloudshell:~/cloud-infra/Project (genuine-grid-327615)$ docker tag adityadw/sonarqube_sonarscanner:latest gcr.io/genuine-grid-327615/adityadw/sonarqube_sonarscanner:1
adidwivedi96@cloudshell:~/cloud-infra/Project (genuine-grid-327615)$ docker push gcr.io/genuine-grid-327615/adityadw/sonarqube_sonarscanner:1
The push refers to repository [gcr.io/genuine-grid-327615/adityadw/sonarqube_sonarscanner]
6cbb6713d8e: Layer already exists
e3e2f421d4ed: Layer already exists
93d8874a9cc2: Layer already exists
e2ab06d8af92: Layer already exists
1: digest: sha256:76354b0d878124ed3a9160874a078cfdc2020144dea399c32a40d7a2d0acb3e5 size: 1159
adidwivedi96@cloudshell:~/cloud-infra/Project (genuine-grid-327615)$
```

4. Create the sonarqube deployment as follows, exposing internal port 9000

```
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: sonar
5    labels:
6      app: sonar
7  spec:
8    selector:
9      matchLabels:
10       app: sonar
11    replicas: 2
12    minReadySeconds: 15
13    strategy:
14      type: RollingUpdate
15      rollingUpdate:
16        maxUnavailable: 1
17        maxSurge: 1
18    template:
19      metadata:
20        labels:
21          app: sonar
22      spec:
23        containers:
24          - image: adityadw/sonarqube_sonarscanner
25            imagePullPolicy: Always
26            name: sonar
27            ports:
28              - containerPort: 9000
```

5. `kubectl apply -f sonarqube_deployment.yaml`
6. Create sonarqube service as follows, mapping internal port 9000 to external port 9000. It's of type LoadBalancer so that it can be accessed via public IP on browser for testing

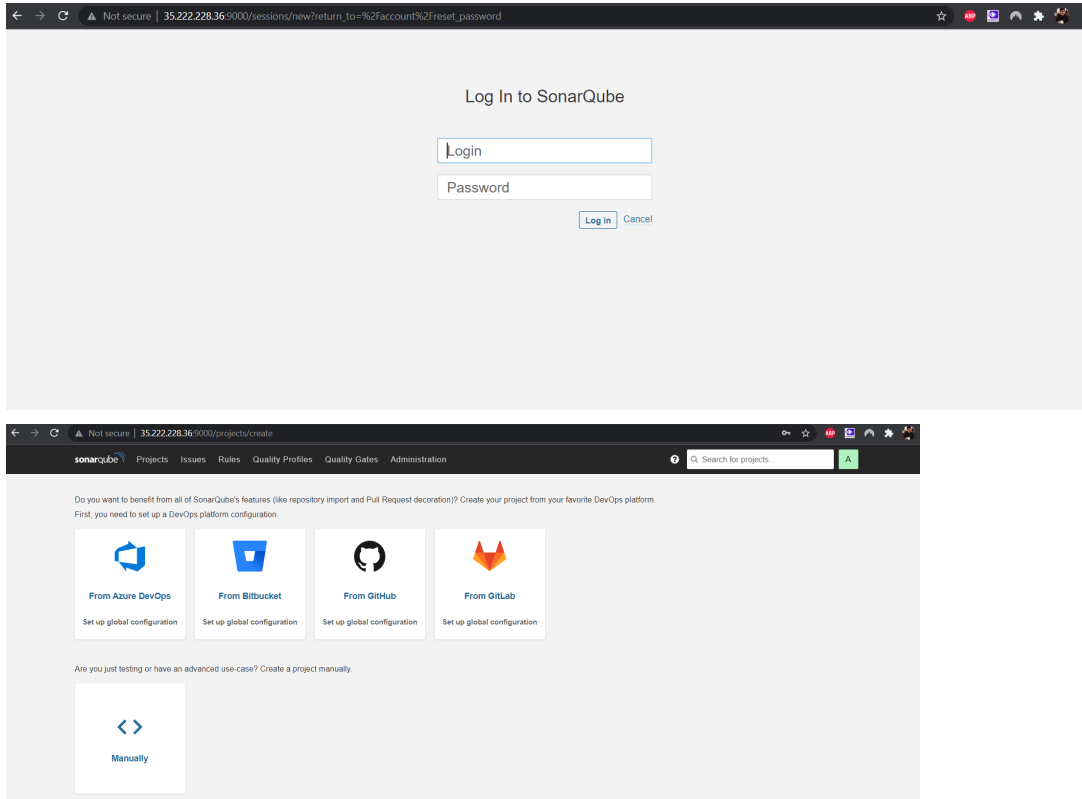
```
sonarqube_service.yaml > ...
1  apiVersion: v1
2  kind: Service
3  metadata:
4    name: sonar
5  spec:
6    type: LoadBalancer
7    ports:
8      - port: 9000
9        protocol: TCP
10       targetPort: 9000
11    selector:
12      app: sonar
```

7. `kubectl apply -f sonarqube_service.yaml`



```
adidwivedi96@cloudshell:~/clud-infra/Project (genuine-grid-327615)$ kubectl apply -f sonarqube_deployment.yaml
deployment.apps/sonar created
adidwivedi96@cloudshell:~/clud-infra/Project (genuine-grid-327615)$ kubectl apply -f sonarqube_service.yaml
service/sonar created
adidwivedi96@cloudshell:~/clud-infra/Project (genuine-grid-327615)$
```

8. We get the following output when we visit the sonarqube



9. We can run sonar-scanner commands within the pod and get results as seen -

```
bash-5.1# sonar-scanner -D sonar.password='helloworld'
INFO: Scanner configuration file: /opt/sonarqube/sonar-scanner-4.6.2.2472/conf/sonar-scanner.properties
INFO: Project root configuration file: NONE
INFO: SonarScanner 4.6.2.2472
INFO: Java 11.0.11 Alpine (64-bit)
INFO: Linux 5.4.144+ amd64
INFO: User cache: /root/.sonar/cache
INFO: Scanner configuration file: /opt/sonarqube/sonar-scanner-4.6.2.2472/conf/sonar-scanner.properties
INFO: Project root configuration file: NONE
INFO: Analyzing on SonarQube server 9.1.0
INFO: Default locale: "en_US", source code encoding: "UTF-8"
INFO: Load global settings
INFO: Load global settings (done) | time=338ms
INFO: Server id: BF41A1F2-AXzF2NSm20tCr7wvp_on
INFO: User cache: /root/.sonar/cache
INFO: Load/download plugins
INFO: Load plugins index
INFO: Load plugins index (done) | time=255ms
INFO: Load/download plugins (done) | time=6196ms
INFO: Process project properties
INFO: Process project properties (done) | time=2ms
INFO: Execute project builders
INFO: Execute project builders (done) | time=4ms
INFO: Project key: myProject
INFO: Base dir: /opt/sonarqube
INFO: Working dir: /opt/sonarqube/.scannerwork
INFO: Load project settings for component key: 'myProject'
INFO: Load quality profiles
INFO: Load quality profiles (done) | time=677ms
INFO: Load active rules
INFO: Load active rules (done) | time=8233ms
WARN: SCM provider autodetection failed. Please use "sonar.scm.provider" to define SCM of your project, or disable the SCM Sensor in the project settings.
INFO: Indexing files...
INFO: Project configuration:
INFO: Load project repositories
```

10. To do so the commands used are -

- a. `kubectl exec --stdin --tty sonar-56b9b86486-t82gv -- /bin/bash`
- b. `sonar-scanner -D sonar.password='password'`

## Terminal

1. `docker pull adityadw/terminal`
2. `docker tag adityadw/terminal:latest gcr.io/genuine-grid-327615/adityadw/terminal:1`
3. `docker push gcr.io/genuine-grid-327615/adityadw/terminal:1`

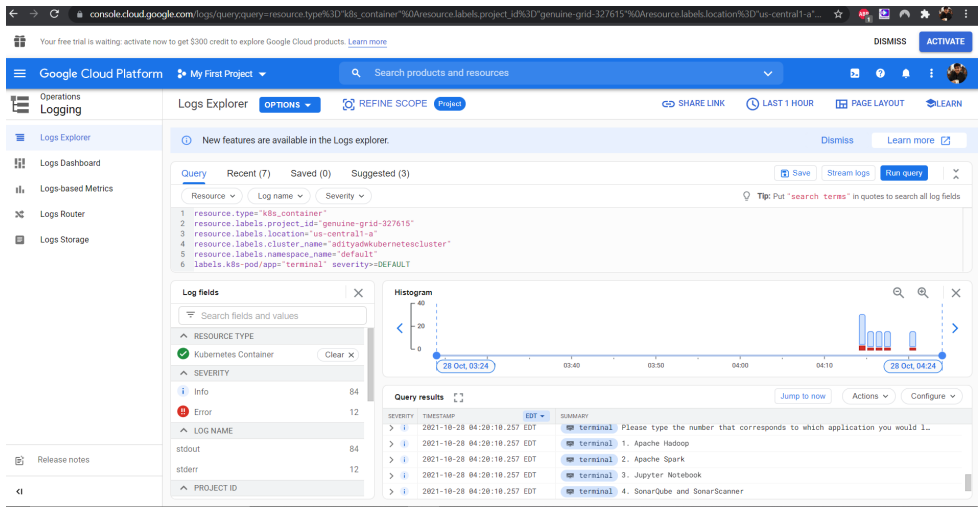
```
adidwivedi96@cloudshell:~/cloud-infra/Project (genuine-grid-327615)$ docker pull adityadw/terminal
Using default tag: latest
latest: Pulling from adityadw/terminal
bb7d5a84853b: Pull complete
f02b617c6a8c: Pull complete
d32e17419b7e: Pull complete
c9d2d81226a4: Pull complete
3c24ae8b6604: Pull complete
8a4322d1621d: Pull complete
b777982287b6: Pull complete
2c5fb32d4bef: Pull complete
4f3be23cccd3: Pull complete
f4adb63687b: Pull complete
4f4fb700ef54: Pull complete
Digest: sha256:32811791a3faf810ba7a4fff60627ec65ed16529153d0fa607353d6cfa9d63a7
Status: Downloaded newer image for adityadw/terminal:latest
docker.io/adityadw/terminal:latest
adidwivedi96@cloudshell:~/cloud-infra/Project (genuine-grid-327615)$ docker tag adityadw/terminal:latest gcr.io/genuine-grid-327615/adityadw/terminal:1
adidwivedi96@cloudshell:~/cloud-infra/Project (genuine-grid-327615)$ docker push gcr.io/genuine-grid-327615/adityadw/terminal:1
The push refers to repository [gcr.io/genuine-grid-327615/adityadw/terminal]
5f70bf18a086: Layer already exists
402eb203ea7c: Pushed
636159d907db: Layer already exists
485115e86f04: Layer already exists
30f3545290f2: Layer already exists
```

4. Create the following deployment yaml

```
terminal_deployment.yaml > {} spec > {} template > {} s|
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: terminal
5    labels:
6      app: terminal
7  spec:
8    selector:
9      matchLabels:
10       app: terminal
11    replicas: 2
12    minReadySeconds: 15
13    strategy:
14      type: RollingUpdate
15      rollingUpdate:
16        maxUnavailable: 1
17        maxSurge: 1
18    template:
19      metadata:
20        labels:
21          app: terminal
22      spec:
23        containers:
24          - image: adityadw/terminal
25            imagePullPolicy: Always
26            name: terminal
27
```

The logs show the following output -

2021-10-28 04:20:22.314 EDT	terminal	Please type the number that corresponds to which application you would l...
2021-10-28 04:20:22.314 EDT	terminal	1. Apache Hadoop
2021-10-28 04:20:22.314 EDT	terminal	2. Apache Spark
2021-10-28 04:20:22.314 EDT	terminal	3. Jupyter Notebook
2021-10-28 04:20:22.314 EDT	terminal	4. SonarQube and SonarScanner
2021-10-28 04:20:22.318 EDT	terminal	Type a number here >



Deployment output

Google Cloud Platform		My First Project		Search products and resources	
Kubernetes Engine		Workloads		REFRESH DEPLOY DELETE	
Clusters		Cluster Namespace		RESET SAVE	
Workloads		Workloads are deployable units of computing that can be created and managed in a cluster.		OVERVIEW COST OPTIMISATION PREVIEW	
Services & Ingress		Filter		Is system object: False Filter workloads	
Applications		Name		Status Type Pods Namespace Cluster	
Configuration		jupyter		OK Deployment 2/2 default adityadwkubernetescluster	
Storage		master-hadoop		OK Deployment 2/2 default adityadwkubernetescluster	
Object Browser		sonar		OK Deployment 2/2 default adityadwkubernetescluster	
Migrate to containers		spark		OK Deployment 2/2 default adityadwkubernetescluster	
Config Management		terminal		Does not have minimum availability Deployment 2/2 default adityadwkubernetescluster	
		worker-hadoop		OK Deployment 2/2 default adityadwkubernetescluster	
		worker-hadoop-2		OK Deployment 2/2 default adityadwkubernetescluster	

```
adidwivedi96@cloudshell:~/clud-infra/Project (genuine-grid-327615)$ kubectl get svc
NAME                TYPE          CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
hadoop-master       LoadBalancer  10.12.6.234     34.136.101.201   9870:31998/TCP   102m
jupyter             LoadBalancer  10.12.7.53      34.132.252.237   8888:31175/TCP   173m
kubernetes          ClusterIP     10.12.0.1       <none>           443/TCP          177m
master-hadoop       LoadBalancer  10.12.10.131    34.136.215.85    9870:31891/TCP,9000:30961/TCP 99m
sonar               LoadBalancer  10.12.15.141    35.222.228.36    9000:32661/TCP   56m
spark              LoadBalancer  10.12.6.231     34.135.56.132    8080:30921/TCP   152m
adidwivedi96@cloudshell:~/clud-infra/Project (genuine-grid-327615)$
```

Building docker images

Jupyter Notebook

1. `docker pull mcr.microsoft.com/mmlspark/release`
2. `docker build -t adityadw/jupyter_notebook .` → runs the jupyter notebook command
3. `docker run -it -p 8888:8888 adityadw/jupyter_notebook`

## Apache Spark

1. `docker pull bitnami/spark`  
URL - <https://hub.docker.com/r/bitnami/spark/>

## Apache Hadoop

1. `docker pull bde2020/hadoop-namenode:2.0.0-hadoop3.2.1-java8`
2. `docker pull bde2020/hadoop-datanode:2.0.0-hadoop3.2.1-java8`

URLs - <https://hub.docker.com/r/bde2020/hadoop-namenode>,  
<https://hub.docker.com/r/bde2020/hadoop-datanode>

## SonarQube and Sonarscanner

1. `docker pull sonarqube`
2. Download sonar-scanner for any platform which requires pre-installed JVM from here <https://binaries.sonarsource.com/Distribution/sonar-scanner-cli/sonar-scanner-cli-4.6.2.2472.zip>
3. Edit sonar-scanner configuration file located under config folder as follows -

```
#Configure here general information about the environment, such as SonarQube server connection details for example
#No information about specific project should appear here

#----- Default SonarQube server
sonar.host.url=http://localhost:9000

#----- Default source code encoding
sonar.sourceEncoding=UTF-8

#Default sonarqube connection and project config
sonar.login=admin
sonar.password=admin
sonar.projectKey=myProject
```

4. Copy the sonar-scanner directory to container and set the PATH variable to bin directory of sonar-scanner as follows

```
1 FROM sonarqube
2 COPY . /opt/sonarqube
3 ENV PATH=$PATH:/opt/sonarqube/sonar-scanner-4.6.2.2472/bin/
```

5. Sonar-scanner command works in container as follows -

```

❏ docker exec -it 99731d8386ea80a04e9707327eda66d645685f197b038660b36334751643a348 /bin/sh
/opt/sonarqube # sonar-scanner
INFO: Scanner configuration file: /opt/sonarqube/sonar-scanner-4.6.2.2472/conf/sonar-scanner.properties
INFO: Project root configuration file: NONE
INFO: SonarScanner 4.6.2.2472
INFO: Java 11.0.11 Alpine (64-bit)
INFO: Linux 5.10.16.3-microsoft-standard-WSL2 amd64
INFO: User cache: /root/.sonar/cache
INFO: Scanner configuration file: /opt/sonarqube/sonar-scanner-4.6.2.2472/conf/sonar-scanner.properties
INFO: Project root configuration file: NONE
INFO: Analyzing on SonarQube server 9.1.0
INFO: Default locale: "en_US", source code encoding: "UTF-8"
INFO: Load global settings
INFO: Load global settings (done) | time=398ms
INFO: Server id: BF41A1F2-AXzFDRLqwjnl2POX-sxL
INFO: User cache: /root/.sonar/cache
INFO: Load/download plugins
INFO: Load plugins index
INFO: Load plugins index (done) | time=319ms
INFO: Load/download plugins (done) | time=5299ms
INFO: Process project properties
INFO: Process project properties (done) | time=1ms
INFO: Execute project builders
INFO: Execute project builders (done) | time=3ms
INFO: Project key: myProject
INFO: Base dir: /opt/sonarqube
INFO: Working dir: /opt/sonarqube/.scannerwork
INFO: Load project settings for component key: 'myProject'
INFO: Load quality profiles
INFO: Load quality profiles (done) | time=456ms
INFO: Load active rules
INFO: Load active rules (done) | time=5978ms
WARN: SCM provider autodetection failed. Please use "sonar.scm.provider" to define SCM of your project, or disable the SCM Sensor in the project settings.
INFO: Indexing files...
INFO: Project configuration:
INFO: Load project repositories
INFO: Load project repositories (done) | time=183ms
INFO: 1003 files indexed
INFO: Quality profile for js: Sonar way
INFO: Quality profile for web: Sonar way
INFO: Quality profile for xml: Sonar way
INFO: ----- Run sensors on module myProject
INFO: Load metrics repository

```

6. docker build -t adityadw/sonarqube\_sonarscanner .

7. docker push adityadw/sonarqube\_sonarscanner

## Terminal

Created python program to simulate the terminal

Created a dockerfile from python base image to run the script

docker build -t adityadw/terminal .

docker push adityadw/terminal

#references - [https://hub.docker.com/\\_/microsoft-mmlspark-release](https://hub.docker.com/_/microsoft-mmlspark-release),  
<https://github.com/big-data-europe/docker-hadoop/blob/master/docker-compose.yml>,  
<https://www.datamechanics.co/blog-post/optimized-spark-docker-images-now-available>,  
<https://computingforgeeks.com/how-to-install-apache-spark-on-ubuntu-debian/>,  
<https://docs.sonarqube.org/latest/analysis/scan/sonarscanner/>  
<https://www.geeksforgeeks.org/python-3-input-function/>  
<https://stackoverflow.com/questions/27093612/in-a-dockerfile-how-to-update-path-environment-variable>  
<https://docs.sonarqube.org/latest/analysis/scan/sonarscanner/>  
<https://www.edureka.co/community/7415/copy-is-not-working-in-docker>  
<https://github.com/mohamedfarag/14-848-extra-credit-project>  
<https://github.com/rinormaloku/k8s-mastery/tree/master/resource-manifests>  
<https://kubernetes.io/docs/tasks/debug-application-cluster/get-shell-running-container/>

